

GET A ROOM : ML HACKATHON

Identify the habitability score of a property

Problem Statement

Finding the correct property to live in is a crucial task while moving to a new city/location. An inappropriate property can make our life miserable. Can AI help us find better places?

Task

You have given a relevant dataset about various properties in the USA. Your task is to identify the habitability score of the property.

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
```

Warning: you are connected to a GPU runtime, but not utilizing the GPU. [Change to a standard runtime](#) ✕

```
1 print(train.shape)
2 train.head(2)
```

(39499, 15)

	Property_ID	Property_Type	Property_Area	Number_of_Windows	Number_of_Doors	Fu
0	0x21e3	Apartment	106	NaN	1	Semi_
1	0x68d4	Apartment	733	2.0	2	Ur



```
1 train.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39499 entries, 0 to 39498
Data columns (total 15 columns):
```

```
#      Column      Non-Null Count  Dtype
---  -
0      Property_ID    39499 non-null  object
1      Property_Type   39499 non-null  object
2      Property_Area    39499 non-null  int64
3      Number_of_Windows 37845 non-null  float64
4      Number_of_Doors   39499 non-null  int64
5      Furnishing        38457 non-null  object
6      Frequency_of_Powercuts 38116 non-null  float64
7      Power_Backup       39499 non-null  object
8      Water_Supply       39499 non-null  object
9      Traffic_Density_Score 39499 non-null  float64
10     Crime_Rate         38712 non-null  object
11     Dust_and_Noise     38280 non-null  object
12     Air_Quality_Index   39499 non-null  float64
13     Neighborhood_Review 39499 non-null  float64
14     Habitability_score  39499 non-null  float64
dtypes: float64(6), int64(2), object(7)
memory usage: 4.5+ MB
```

```
1 test = pd.read_csv("/content/test.csv")
```

```
1 test.shape

(10500, 14)
```

Warning: you are connected to a GPU runtime, but not utilizing the GPU. [Change to a standard runtime](#) 

```
Property_ID      0
Property_Type    0
Property_Area     0
Number_of_Windows 1654
Number_of_Doors   0
Furnishing       1042
Frequency_of_Powercuts 1383
Power_Backup      0
Water_Supply      0
Traffic_Density_Score 0
Crime_Rate        787
Dust_and_Noise    1219
Air_Quality_Index  0
Neighborhood_Review 0
Habitability_score 0
dtype: int64
Property_ID      0
Property_Type    0
Property_Area     0
Number_of_Windows 445
Number_of_Doors   0
Furnishing       257
Frequency_of_Powercuts 366
Power_Backup      0
```

```
Water_Supply      0
Traffic_Density_Score  0
Crime_Rate        212
Dust_and_Noise    330
Air_Quality_Index  0
Neighborhood_Review  0
dtype: int64
```

```
1 from sklearn.impute import SimpleImputer, KNNImputer
```

```
1 impute = SimpleImputer(strategy = "most_frequent")
2 knnimputer = KNNImputer(n_neighbors=10)
```

```
1 y = train.iloc[:, -1]
2 train.drop(train.columns[-1], axis=1, inplace=True)
```

```
1 train2 = train.iloc[:, [3, 6]]
2 train2 = knnimputer.fit_transform(train2)
3 train2 = pd.DataFrame(train2)
```

```
1 train3 = train.iloc[:, [5, 10, 11]]
2 train3 = impute.fit_transform(train3)
3 train3 = pd.DataFrame(train3)
```

Warning: you are connected to a GPU runtime, but not utilizing the GPU. [Change to a standard runtime](#) ✕

```
1 train.iloc[:, 3] = train2.iloc[:, 0]
2 train.iloc[:, 6] = train2.iloc[:, 1]
3 train.iloc[:, 5] = train3.iloc[:, 0]
4 train.iloc[:, 10] = train3.iloc[:, 1]
5 train.iloc[:, 11] = train3.iloc[:, 2]
```

```
1 test2 = test.iloc[:, [3, 6]]
2 test2 = knnimputer.fit_transform(test2)
3 test2 = pd.DataFrame(test2)
```

```
1 test3 = test.iloc[:, [5, 10, 11]]
2 test3 = impute.fit_transform(test3)
3 test3 = pd.DataFrame(test3)
```

```
1 test.iloc[:, 3] = test2.iloc[:, 0]
2 test.iloc[:, 6] = test2.iloc[:, 1]
3 test.iloc[:, 5] = test3.iloc[:, 0]
4 test.iloc[:, 10] = test3.iloc[:, 1]
5 test.iloc[:, 11] = test3.iloc[:, 2]
```

```
1 from sklearn.preprocessing import OneHotEncoder
```

```

2 from sklearn.preprocessing import StandardScaler
3 ohe = OneHotEncoder()
4 sc = StandardScaler()

1 train2 = train.iloc[:,[1,5,7,8,10,11]]
2 train2 = ohe.fit_transform(train2)
3 train2 = pd.DataFrame(train2.toarray())
4
5 train3 = train.iloc[:,[2,3,4,6,9,12,13]]
6 train3 = sc.fit_transform(train3)
7 train3 = pd.DataFrame(train3)
8
9 train = pd.concat([train2,train3],axis=1)

1 test2 = test.iloc[:,[1,5,7,8,10,11]]
2 test2 = ohe.fit_transform(test2)
3 test2 = pd.DataFrame(test2.toarray())
4
5 test3 = test.iloc[:,[2,3,4,6,9,12,13]]
6 test3 = sc.fit_transform(test3)
7 test3 = pd.DataFrame(test3)
8
9 test = pd.concat([test2,test3],axis=1)

```

Warning: you are connected to a GPU runtime, but not utilizing the GPU.

[Change to a standard runtime](#) ✕

```

3 from sklearn.tree import DecisionTreeRegressor
4 from sklearn.ensemble import RandomForestRegressor
5 from sklearn.svm import SVR
6
7 lr = LinearRegression()
8 lr.fit(x_train,y_train)
9 y_pred = lr.predict(x_test)
10
11 dtr = DecisionTreeRegressor()
12 dtr.fit(x_train,y_train)
13 y_pred2 = dtr.predict(x_test)
14
15 rfr = RandomForestRegressor()
16 rfr.fit(x_train,y_train)
17 y_pred3 = rfr.predict(x_test)
18
19 svr = SVR()
20 svr.fit(x_train,y_train)
21 y_pred4 = svr.predict(x_test)
22
23 return y_pred,y_pred2,y_pred3,y_pred4

```

```

1 def regressor2(x_train,x_test,y_train,y_test):

```

```

2  from sklearn.metrics import mean_absolute_error
3  from sklearn.metrics import r2_score
4  from sklearn.linear_model import LinearRegression
5  from sklearn.tree import DecisionTreeRegressor
6  from sklearn.ensemble import RandomForestRegressor
7  from sklearn.svm import SVR
8
9  lr = LinearRegression()
10 lr.fit(x_train,y_train)
11 y_pred = lr.predict(x_test)
12 lre = mean_absolute_error(y_test,y_pred)
13 print("Linear Regression\nmean absolute error : ",lre,"\n")
14 print("R2_score : ",r2_score(y_test,y_pred),"\n")
15
16 dtr = DecisionTreeRegressor()
17 dtr.fit(x_train,y_train)
18 y_pred2 = dtr.predict(x_test)
19 dtre = mean_absolute_error(y_test,y_pred2)
20 print("Decision Tree Classifier\nmean absolute error : ",dtre,"\n")
21 print("R2_score : ",r2_score(y_test,y_pred2),"\n")
22
23 rfr = RandomForestRegressor()
24 rfr.fit(x_train,y_train)
25 y_pred3 = rfr.predict(x_test)
26 rfrc = mean_absolute_error(y_test,y_pred3)
27 print("Random Forest Regressor\nmean absolute error : ",rfrc,"\n")
28
29
30 svr = SVR()
31 svr.fit(x_train,y_train)
32 y_pred4 = svr.predict(x_test)
33 svrc = mean_absolute_error(y_test,y_pred4)
34 print("Support Vector Regression\nmean absolute error : ",svrc,"\n")
35 print("R2_score : ",r2_score(y_test,y_pred4),"\n")
36
37 return y_pred,y_pred2,y_pred3,y_pred4

```

Warning: you are connected to a GPU runtime, but not utilizing the GPU.

[Change to a standard runtime](#) ✕

```
1 p1,p2,p3,p4 = regressor1(train,test,y)
```

```
1 from sklearn.metrics import mean_absolute_error
2 mean_absolute_error(p1,p2)
```

```
7.393517008928571
```

```
1 mean_absolute_error(p2,p3)
```

```
4.156624790476189
```

```
1 mean_absolute_error(p3,p4)
```

```
3.58908326486187
```

```
1 mean_absolute_error(p4,p2)
```

```
5.812582066087356
```

```
1 p3
```

```
array([30.5028, 80.1084, 67.392 , ..., 75.1799, 79.6257, 79.4099])
```

```
1 test = pd.read_csv("/content/test.csv")
```

```
2 id = test.iloc[:,0]
```

```
1 submission = pd.DataFrame({"Property_ID":id,"Habitability_score":p3})
```

```
1 submission
```

	Property_ID	Habitability_score
0	0x6e93	30.5028
1	0x8787	80.1084
3	0x9dbd	73.0809
4	0xbfde	75.4409
...
10495	0x423d	64.4912
10496	0x78c5	81.0152
10497	0xbf3	75.1799
10498	0x305b	79.6257
10499	0x5cff	79.4099


10500 rows × 2 columns

```
1 submission.to_csv("Submission.csv",index=False)
```

```
1 pd.read_csv("Submission.csv")
```

Warning: you are connected to a GPU runtime, but not utilizing the GPU.

[Change to a standard runtime](#) ×

	Property_ID	Habitability_score	
0	0x6e93	30.5028	
1	0x8787	80.1084	
2	0x6c17	67.3920	
3	0x9dbd	73.0809	
4	0xbfde	75.4409	
...	
10495	0x423d	64.4912	
10496	0x78c5	81.0152	
10497	0xbf3	75.1799	
10498	0x305b	79.6257	
10499	0x5cff	79.4099	

1

Warning: you are connected to a GPU runtime, but not utilizing the GPU.

[Change to a standard runtime](#)

×